

Assignment -4

Mall Customers

Assignment Date	24 September 2022
Student Name	Paul Jabez Talakayala
Student Roll Number	195002079
Maximum Marks	2 Marks

Load the dataset into the tool.

```
[29] import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

[30] df=pd.read_csv("Mall_Customers.csv")
df.head()
```

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

1.Perform Below Visualizations. Univariate Analysis Bivariate Analysis Multivariate Analysis

```
#Univariate Analysis
#Age
plt.hist(df['Age'])
plt.show()

#Spending Score
plt.hist(df['Spending Score (1-100)'])
plt.show()

#Annual Income
plt.hist(df['Annual Income (k$)'])
plt.show()

#Bi-Variate Analysis
#Age vs Spending Score
plt.scatter(df['Age'],df['Spending Score (1-100)'])
plt.show()

#Age vs Annual Income
plt.scatter(df['Age'],df['Annual Income (k$)'])
plt.show()
```

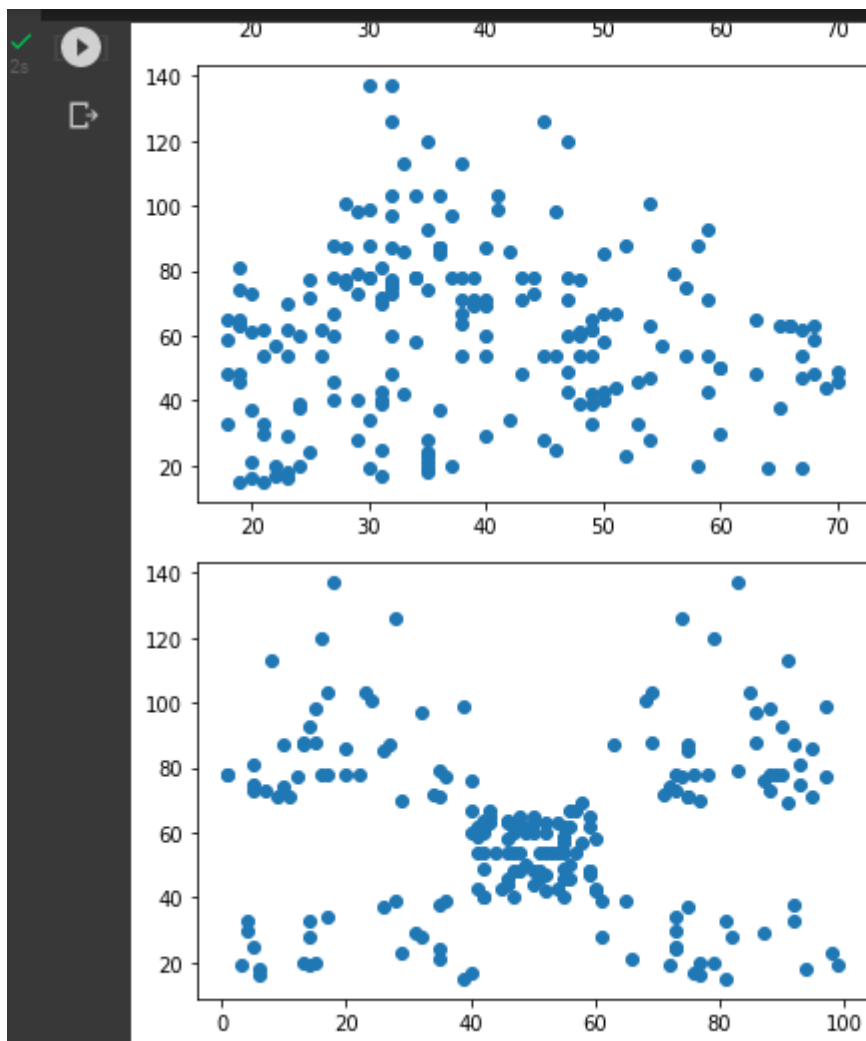
```

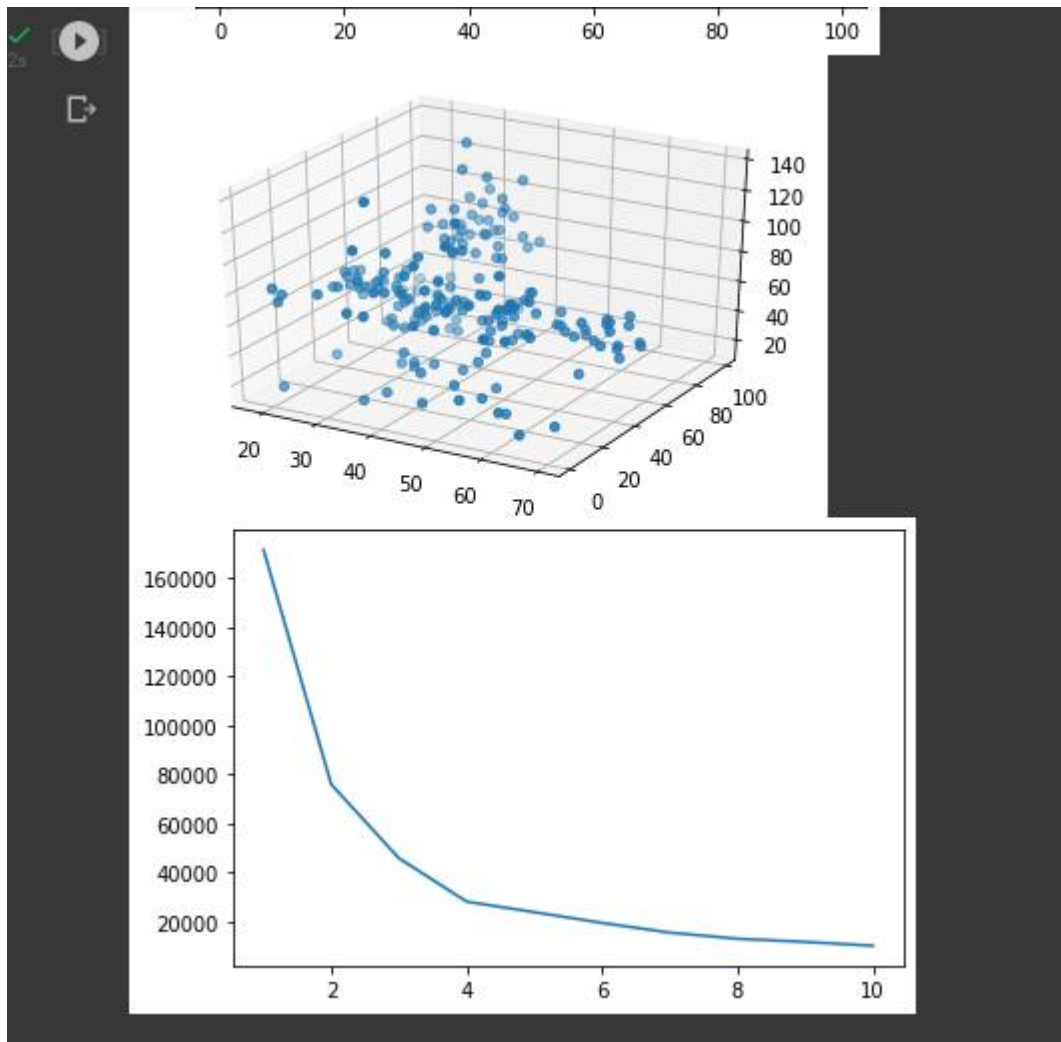
#Spending Score vs Annual Income
plt.scatter(df['Spending Score (1-100)'],df['Annual Income (k$)'])
plt.show()

#Multi-Variate Analysis
#Age vs Spending Score vs Annual Income
from mpl_toolkits.mplot3d import Axes3D
fig=plt.figure()
ax=fig.add_subplot(111,projection='3d')
ax.scatter(df['Age'],df['Spending Score (1-100)'],df['Annual Income (k$)'])
plt.show()

#K-Means Clustering
#Age vs Spending Score
X=df.iloc[:,[2,4]].values
from sklearn.cluster import KMeans
wcss=[]
for i in range(1,11):
    kmeans=KMeans(n_clusters=i,init='k-means++',random_state=42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)
plt.plot(range(1,11),wcss)
plt.show()

```





▼ Perform descriptive statistics on the dataset.

✓ [32] df.info()

0s

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                           200 non-null    int64
1   Genre                                200 non-null    object
2   Age                                  200 non-null    int64
3   Annual Income (k$)                   200 non-null    int64
4   Spending Score (1-100)                200 non-null    int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

✓ [33] df.isnull().sum()

0s

```
CustomerID      0
Genre            0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

✓ [34] df.columns

0s

```
Index(['CustomerID', 'Genre', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)'],
      dtype='object')
```

✓ [35] df.shape

0s

```
(200, 5)
```

✓ [36] df.dtypes

0s

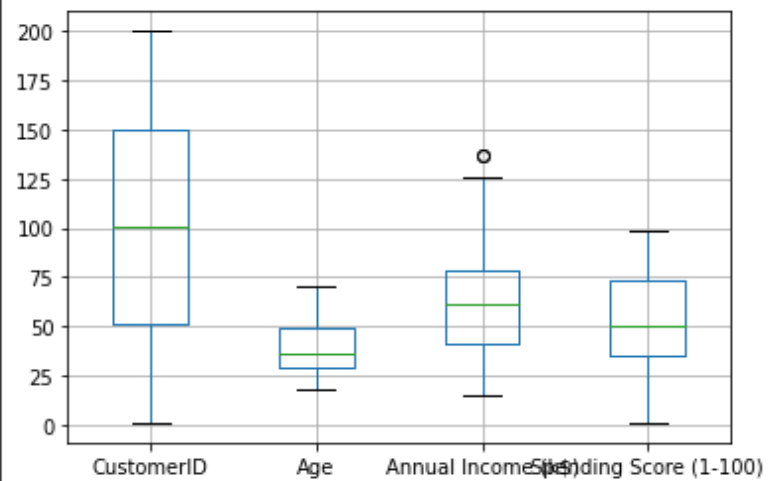
```
CustomerID      int64
Genre            object
Age             int64
Annual Income (k$)  int64
Spending Score (1-100)  int64
dtype: object
```

Find the outliers and replace them outliers

✓ [37] df.boxplot()

0s

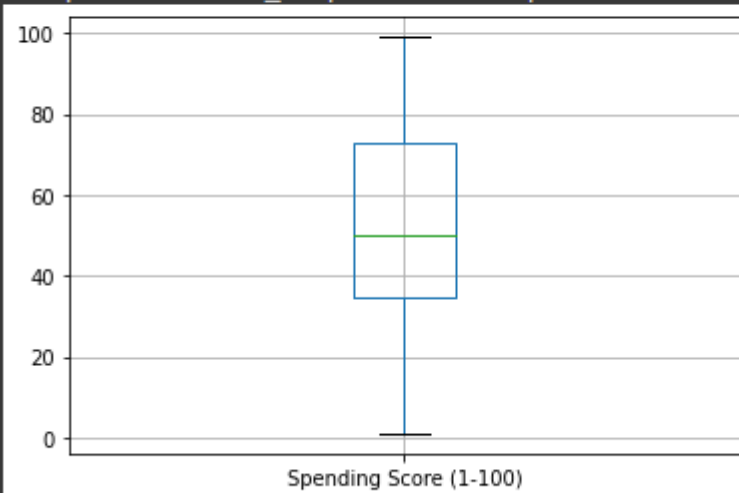
<matplotlib.axes._subplots.AxesSubplot at 0x7fe64dc80510>



✓ [38] df.boxplot(column='Spending Score (1-100)')

0s

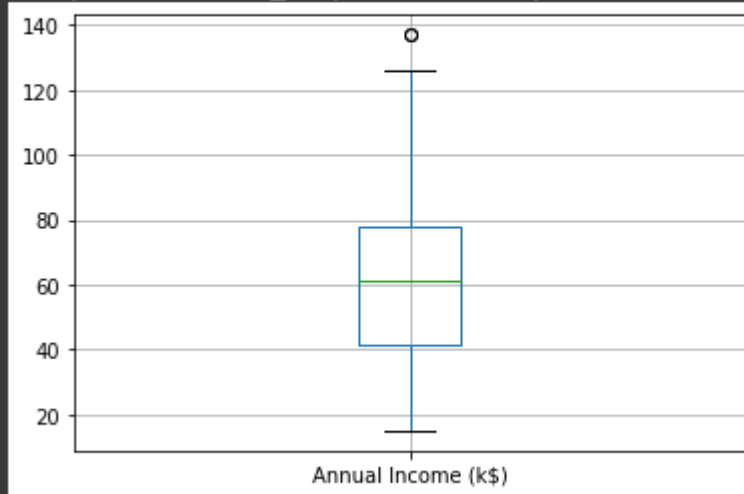
<matplotlib.axes._subplots.AxesSubplot at 0x7fe64d53b990>



✓
0s

```
[39] df.boxplot(column='Annual Income (k$)')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fe64d4b4f50>
```



✓
0s

```
[40] df.boxplot(column='Age')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fe64d42cad0>
```



▼ Check for Categorical columns and perform encoding.

```
[41] df.select_dtypes(include='object').columns
```

```
Index(['Genre'], dtype='object')
```

```
[42] df=pd.get_dummies(df,columns=['Genre'])  
df.head()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)	Genre_Female	Genre_Male
0	1	19	15	39	0	1
1	2	21	15	81	0	1
2	3	20	16	6	1	0
3	4	23	16	77	1	0
4	5	31	17	40	1	0

▼ Scaling the data

```
[43] from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()  
df_scaled=scaler.fit_transform(df)  
df_scaled
```

```
array([[ -1.7234121, -1.42456879, -1.73899919, -0.43480148, -1.12815215,  
         1.12815215],  
       [-1.70609137, -1.28103541, -1.73899919,  1.19570407, -1.12815215,  
         1.12815215],  
       [-1.68877065, -1.3528021 , -1.70082976, -1.71591298,  0.88640526,  
        -0.88640526],  
       ...,  
       [ 1.68877065, -0.49160182,  2.49780745,  0.92395314, -1.12815215,  
         1.12815215],  
       [ 1.70609137, -0.49160182,  2.91767117, -1.25005425, -1.12815215,  
         1.12815215],  
       [ 1.7234121 , -0.6351352 ,  2.91767117,  1.27334719, -1.12815215,  
         1.12815215]])
```

▼ Perform any of the clustering algorithms

```
[44] from sklearn.cluster import KMeans  
kmeans=KMeans(n_clusters=5,random_state=42)  
kmeans.fit(df_scaled)
```

```
KMeans(n_clusters=5, random_state=42)
```

```
kmeans.labels_
```

```
array([4, 4, 2, 2, 2, 2, 2, 2, 3, 2, 3, 2, 2, 2, 3, 4, 2, 4, 3, 2, 4, 4,  
       2, 4, 2, 4, 2, 4, 2, 2, 3, 2, 3, 4, 2, 2, 2, 2, 2, 2, 2, 4, 3, 2,  
       2, 2, 2, 2, 2, 2, 4, 2, 3, 2, 3, 2, 3, 2, 3, 3, 4, 2, 2, 3, 4,  
       2, 2, 4, 2, 3, 2, 2, 2, 3, 4, 2, 3, 2, 2, 3, 4, 3, 2, 2, 3, 2, 2,  
       2, 2, 2, 4, 3, 2, 2, 4, 2, 1, 3, 4, 1, 2, 3, 4, 3, 1, 2, 3, 3, 3,  
       3, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 3, 0, 0, 0,  
       1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1,  
       1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1,  
       0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0,  
       0, 0], dtype=int32)
```

▼ Add the cluster data with the primary dataset

```
[47] df['Cluster']=kmeans.labels_  
df.head()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)	Genre_Female	Genre_Male	Cluster
0	1	19	15	39	0	1	4
1	2	21	15	81	0	1	4
2	3	20	16	6	1	0	2
3	4	23	16	77	1	0	2
4	5	31	17	40	1	0	2

```
kmeans.cluster_centers_
```

```
array([[ 1.00571233, -0.25605987,  0.9485201 , -0.09735253, -1.12815215,  
        1.12815215],  
       [ 0.88269077, -0.19763442,  0.8088102 ,  0.11840576,  0.88640526,  
       -0.88640526],  
       [-0.85997398,  0.07057058, -0.79430582, -0.00647026,  0.88640526,  
       -0.88640526],  
       [-0.53075649,  1.33075947, -0.48486081, -0.42786906, -1.12815215,  
        1.12815215],  
       [-0.88871813, -1.01105596, -0.84837918,  0.47658087, -1.12815215,  
        1.12815215]])
```

▼ Split the data into dependent and independent variables.

```
[48] X=df.drop('Cluster',axis=1)  
y=df['Cluster']
```

▼ Split the data into dependent and independent variables.

```
[49] from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

▼ Build the Model

```
[50] from sklearn.linear_model import LogisticRegression  
model=LogisticRegression()  
model.fit(X_train,y_train)
```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
LogisticRegression())

▼ Train the Model

```
✓ [51] model.score(X_train,y_train)
```

0s

0.95

▼ Test the Model

```
✓ [52] model.score(X_test,y_test)
```

0s

0.95

▼ Measure the performance using Evaluation Metrics.

```
✓ [53] from sklearn.metrics import confusion_matrix,classification_report  
y_pred=model.predict(X_test)  
confusion_matrix(y_test,y_pred)
```

0s

```
array([[ 7,  0,  0,  0,  0],  
       [ 0, 10,  0,  0,  0],  
       [ 0,  0, 11,  0,  0],  
       [ 1,  0,  0,  6,  0],  
       [ 1,  0,  0,  0,  4]])
```

```
✓ [54] print(classification_report(y_test,y_pred))
```

0s

	precision	recall	f1-score	support
0	0.78	1.00	0.88	7
1	1.00	1.00	1.00	10
2	1.00	1.00	1.00	11
3	1.00	0.86	0.92	7
4	1.00	0.80	0.89	5
accuracy			0.95	40
macro avg	0.96	0.93	0.94	40
weighted avg	0.96	0.95	0.95	40